

# AI Cheat Sheet for Traders

Stop prompting like a beginner. Start prompting like a quant.

MOMENTUM PHINANCE • APRIL 2026

## 1. The Three Layers of Every Prompt

An AI does not "think." It predicts the next word based on everything you wrote before it. More specific input means more specific output. Your job is to control what it predicts.

**Layer 1: Role** Who the model is. "You are a risk manager." This is not cosmetic. It activates different training patterns.

**Layer 2: Context** The data it needs. Stock price, account size, IV rank, DTE. If you do not provide it, the model guesses.

**Layer 3: Task** The specific output you want, in the exact format you want it. This is where most people fail.

## 2. The Five Rules That Actually Matter

### 🎯 Rule 1: Define the Output Format

The single highest-leverage change you can make. If you define the structure before the model starts writing, it cannot forget to cover a section.

✗ BAD

"Analyze this trade."

✓ GOOD

"Analyze this covered call. Format as: VERDICT, PREMIUM QUALITY (1-10), ASSIGNMENT RISK, KEY RISK, RECOMMENDATION."

### 🧑‍💻 Rule 2: Assign a Role

Different roles produce genuinely different answers to the same question.

"You are a risk manager." → focuses on what can go wrong

"You are a quant trader maximizing returns." → focuses on optimization

"You are a beginner investor." → focuses on clarity and simplicity

### 🔗 Rule 3: Ask for the Work

For anything involving math, greeks, or probability, tell the model to show its reasoning step by step. The intermediate tokens become context for the final answer.

*"Think through this step by step. State your assumptions, derive the formula, then plug in the numbers."*

### 🔗 Rule 4: Constraints Beat Freedom

Counterintuitively, giving the model less room produces better results. Constraints force prioritization.

- *"In 3 sentences or less"* forces it to pick what matters
- *"Without using generic adjectives"* forces specificity
- *"Assume I know options but not Python"* calibrates depth

### 🔗 Rule 5: Temperature Matters

- |           |                                                     |
|-----------|-----------------------------------------------------|
| 0.0 - 0.2 | Math, code, facts. Deterministic. One right answer. |
| 0.3 - 0.4 | Analysis, trade review. Grounded but thoughtful.    |
| 0.5 - 0.7 | Mixed tasks. Some creativity is okay.               |
| 0.8 - 1.0 | Brainstorming, ideation. Maximum variation.         |

*Most AI tools default to 0.7+. For trading analysis, you want 0.2-0.4. Lower temperature = more precise.*

### 3. Plug-and-Play: The Options Setup Review

Copy this into ChatGPT, Claude, or Gemini before entering any premium-selling trade. Fill in the brackets.

```
# Copy this entire block. Fill in the [brackets].
```

**[ROLE]** You are a strict risk manager reviewing an options trade for a small retail account.

**[SETUP]**

Underlying: **[TICKER]** @ \$**[PRICE]**

Strategy: **[CSP / CC / spread]**

Strike(s): \$**[STRIKE]**

Expiration: **[DATE]** (**[N]** DTE)

Premium: \$**[CREDIT or DEBIT]**

Account: \$**[SIZE]**

**[TASK]** Evaluate:

1. Premium quality: is this worth the risk? (annualized %)
2. Assignment probability: realistic assessment
3. Key risk: the ONE thing that kills this trade
4. Regime check: does this fit current conditions?
5. Recommendation: Enter / Wait / Pass

Be direct. Numbers, not adjectives.

### 4. Plug-and-Play: The Strategy Stress-Test

Use this when evaluating a new strategy or reviewing a backtest.

```
# Use this to pressure-test any strategy before risking capital.
```

**[ROLE]** You are a quantitative analyst reviewing a backtesting methodology.

**[STRATEGY]** **[Describe the strategy]**

**[BACKTEST]** **[Describe what was tested and results]**

**[TASK]** Identify:

1. Survivorship bias risks
2. Look-ahead bias risks
3. Overfitting indicators (too many parameters?  
cherry-picked timeframe?)
4. Transaction cost assumptions: realistic?
5. Regime dependency: when does this break?

For each issue: rate severity  
(critical / moderate / minor) and describe the fix.

## 5. Troubleshooting Your AI

If the AI gives you garbage, it is almost always one of these three problems.

### Problem: Generic answers

**Cause:** You gave it a vague ask. "Should I sell a covered call?"

**Fix:** Add context. "Should I sell a CC on 100 shares of NVDA (cost basis \$620) with earnings in 45 days? IV rank is 65%. I want 2% monthly income."

### Problem: 1,000-word essays when you needed a yes/no

**Cause:** No format constraints.

**Fix:** Add a word limit or exact output structure. "Answer in 3 sentences."

### Problem: "This is not financial advice" hedging

**Cause:** No role assigned. The model defaults to cautious.

**Fix:** Give it a persona. "You are my personal trading analyst. Skip disclaimers. Be direct."

## Quick Reference: The Prompt Skeleton

```
[ROLE]    You are a {role} with {expertise}.
[CONTEXT] {background info the model needs}
[TASK]    {specific thing you want}
[FORMAT]  {exact output structure}
[INPUT]   {the content to analyze}
```

*Get all five right and you will get better answers than 99% of people using the same model.*

Built for the Alpha Dossier Community by [Momentum Phinance](#)

Not financial advice. I build tools, not portfolios. Use at your own risk.